



中山大學
SUN YAT-SEN UNIVERSITY



国家超级计算广州中心
NATIONAL SUPERCOMPUTER CENTER IN GUANGZHOU

Compilation Principle 编译原理

第1讲：词法分析

张献伟

xianweiz.github.io

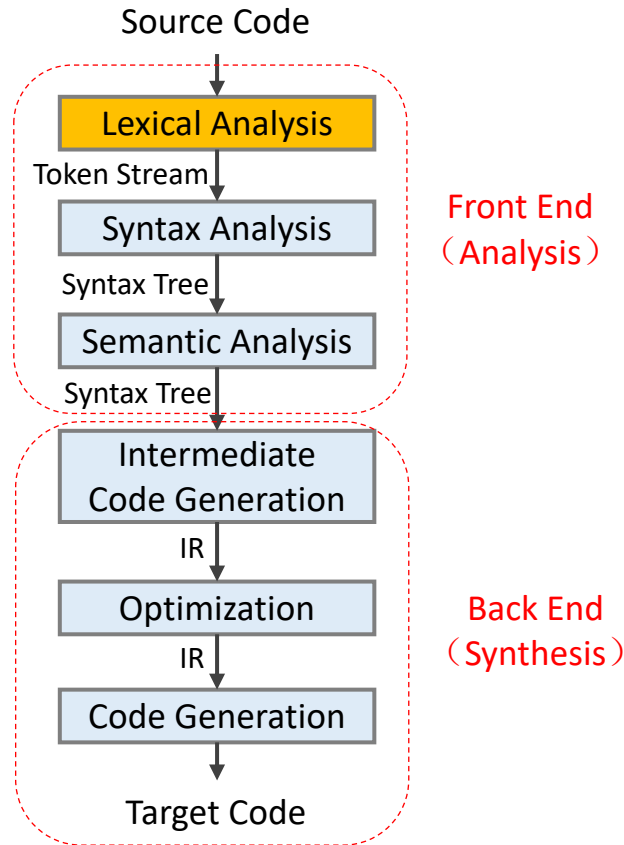
DCS290, Spring 2021



中山大學
SUN YAT-SEN UNIVERSITY



Structure of a Typical Compiler



What is Lexical Analysis[词法分析]?

- Example:


```
/* simple example */  
if (i == j)  
    z = 0;  
else  
    z = 1;
```
- Input: a string of characters
 - “*if (i == j)\n\t\tz = 0; \telse\n\tz = 1; \n*”
- Goal: partition the string into a set of substrings
 - Those substrings are **tokens**
- Steps:
 - Remove comments
 - Identify substrings: ‘if’ ‘(’ ‘i’ ‘==’ ‘j’
 - Identify **token classes**: (keyword, ‘if’), (LPAR, ‘(’), (id, ‘i’)

What is a token[词]?

- **Token**: a “word” in language (smallest unit with meaning)
 - Categorized into classes according to its role in language
 - Token classes in English:
 - Noun, verb, adjective, ...
 - Token classes in a programming language
 - Number, keyword, whitespace, identifier, ...
- Each **token class** corresponds to a set of strings
 - **Numbers**: a non-empty string of digits
 - **Keyword**: a fixed set of reserved words (“for”, “if”, “else”, ...)
 - **Whitespace**: a non-empty sequence of blanks, tabs, newlines
 - **Identifier**: user-defined name of an entity to identify (Q: what are the rules in C language?)

Lexical Analysis: Tokenization[分词]?

- Lexical analysis is also called **Tokenization** (also called Scanner)
 - Partition input string into a sequence of tokens
 - Classify each token according to roles (token class)
 - **Lexeme**: an instance of the corresponding token class, e.g. 'z', '=', '1'
- Pass tokens to syntax analyzer (also called Parser)
 - Parser relies on token classes to identify roles (e.g., a keyword is treated differently than an identifier)



Lexical Analyzer: Design

- Define a finite set of token classes
 - Describe all items of interest
 - Depends on language, design of parser
 - “*if (i == j) \n t \tz = 0; \telse \n \tz = 1; \n*”
 - Keyword, identifier, whitespace, integer
- Label which string belongs to which token class

```
if (i == j)
    z = 0;
else
    z = 1;
```

'==' or '='?

keyword or identifier?

Lexical Analyzer: Implementation

- An implementation must do two things
 - Recognize the token class the substring belongs to
 - Return the value or lexeme of the token
- A token is a tuple (class, lexeme)
- The lexer usually discards “non-interesting” tokens that don’t contribute to parsing, e.g., whitespace, comments
- If token classes are non-ambiguous, tokens can be recognized in a single left-to-right scan of input string
- Problem can occur when classes are ambiguous

Ambiguous Tokens in C++

- C++ template syntax
 - Foo<Bar>
- C++ stream syntax
 - cin >> var

- Ambiguity

- Foo<Bar<Bar>>>
- cin >>> var
- Q: Is '>>>' a stream operator or two consecutive brackets?

```
Template <typename T>
```

```
T getMax(T x, T y) {  
    return (x > y) ? x : y;  
}
```

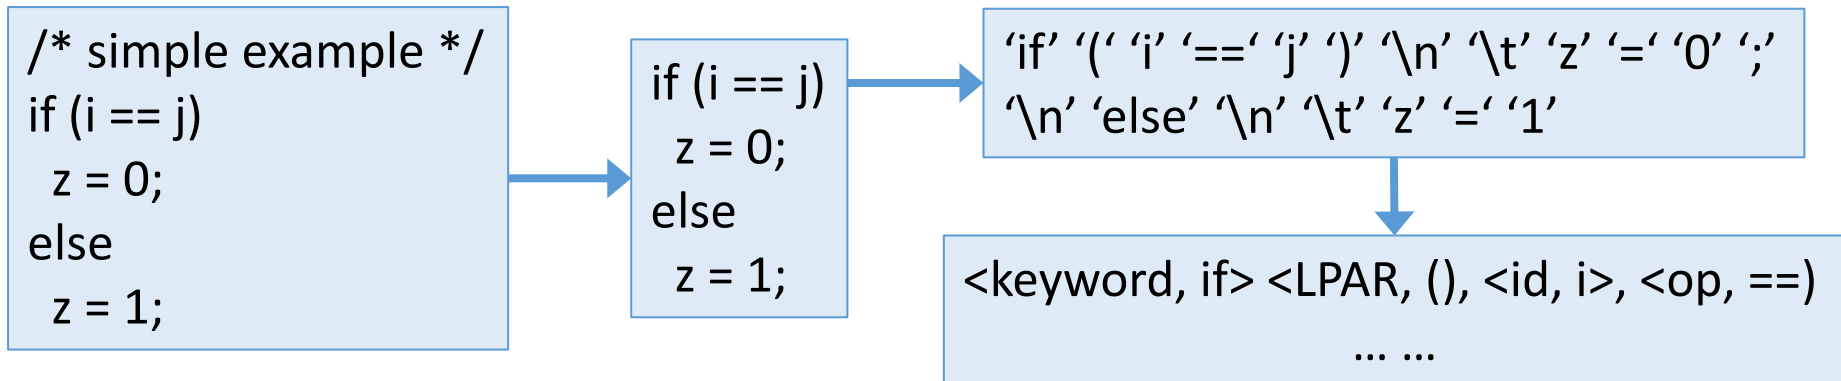
```
int main (int argc, char* argv[]) {  
    getMax<int>(3, 7);  
    getMax<double>(3.0, 2.0);  
    getMax<char>('g', 'e');  
  
    return 0;  
}
```


Look Ahead

- “look ahead” may be required to resolve ambiguity
 - Extracting some tokens requires looking at the larger context or structure
 - Structure emerges only at parsing stage with parse tree
 - Hence, sometimes feedback from parser needed for lexing
 - This complicates the design of lexical analysis
 - Should minimize the amount of look ahead
- Usually tokens do not overlap
 - Tokenizing can be done in one pass w/o parser feedback
 - Clean division between lexical and syntax analyses


Summary: Lexer

- Lexical Analysis
 - Partition the input string to lexeme
 - Identify the token class of each lexeme
- Left-to-right scan => look ahead may be required
 - In reality, lookahead is always needed
 - The amount of lookahead should be minimized



Token Specification

- Recognizing token class: how to describe string patterns
 - i.e., which set of strings belong to which token class?
 - Use regular expressions [正则表达式] to define token class
- Regular Expression is a good way to specify tokens
 - Simple yet powerful (able to express patterns)
 - Tokenizer implementation can be generated automatically from specification (using a translation tool)
 - Resulting implementation is provably efficient



String patterns
describing the class

Language: Definition

- **Alphabet** Σ [字母表]: a finite set of symbols
 - Symbol: letter, digit, punctuation, ...
 - Example: $\{0, 1\}$, $\{a, b, c\}$, ASCII
- **String** [串]: a finite sequence of symbols drawn from Σ
 - Example: aab (length = 3), ϵ (empty string, length = 0)
- **Language** [语言]: a set of strings of the characters drawn from Σ
 - $\Sigma = \{0, 1\}$, then $\{\}$, $\{01, 10\}$, $\{1, 11, 1111, \dots\}$ are all languages over Σ
 - $\{\epsilon\}$ is a language
 - Φ , empty set is also a language

Language: Example

- Examples:
 - Alphabet Σ = (set of) English characters
 - Language L = (set of) English sentences
 - Alphabet Σ = (set of) Digits, +, -
 - Language L = (set of) Integer numbers
- Languages are subsets of all possible strings
 - Not all strings of English characters are sentences
 - Not all sequences of digits and signs are integers

Regular Expressions and Languages

- Need a notion to specify strings in a particular language
 - More complex languages need more complex notations
- **Regular Expression** is a simple notation
 - Can express simple patterns (e.g., repeating sequences)
 - Not powerful enough to express English (or even C)
 - But powerful enough to express tokens (e.g., identifiers)
- Languages that can be expressed using regular expressions are called **Regular Languages**
- More complex languages and expressions will be covered later

Atomic REs[原子]

- Atomic
 - Smallest RE that cannot be broken down further
- **Epsilon or ϵ** character denotes a zero length string
 - $\epsilon = \{""\}$
- **Single character** denotes a set of one string
 - $'c' = \{“c”\}$
- Empty set is $\{ \} = \phi$, not the same as ϵ
 - $\text{Size}(\phi) = 0$
 - $\text{Size}(\epsilon) = 1$
 - $\text{Length}(\epsilon) = 0$

Compound REs[组合]

- **Union[并]**: if A and B are REs, then
$$A|B = \{ s \mid s \in A \text{ or } s \in B \}$$
- **Concatenation[连接]** of sets/strings
$$AB = \{ ab \mid a \in A \text{ and } b \in B \}$$
- **Iteration[迭代]** (Kleene closure)
$$A^* = \bigcup_{i \geq 0} A^i \text{ where } A^i = A \dots A \text{ (i times)}$$

in particular

$$A^* = \{\epsilon\} + A + AA + AAA + \dots$$
$$A^+ = A + AA + AAA + \dots = AA^*$$
- **(A) \equiv A**: A is a RE

RE and RL

- The regular expressions (REs) over Σ are the total set of expressions that can be constructed using components:
 - ϵ
 - $'c'$ where $c \in \Sigma$
 - $A|B$ where A, B are REs over Σ
 - AB where A, B are REs over Σ
 - A^* where A is a RE over Σ
- The regular languages (RLs) over Σ are the total set of languages that can be expressed using REs:
 - $L(\epsilon) = \{''''\}$
 - $L('c') = \{''c''\}$
 - $L(A|B) = L(A) \cup L(B)$
 - $L(AB) = \{ab \mid a \in L(A) \text{ and } b \in L(B)\}$
 - $L(A^*) = \bigcup_{i \geq 0} L(A^i)$

Operator Precedence[优先级]

- RE operator precedence

- (A)
- A^*
- AB
- $A|B$

- Example: $ab^*c|d$

- $a\underline{(b^*)}c|d$
- $\underline{(a(b^*))}c|d$
- $\underline{((a(b^*))c)}|d$

Common REs

- **At least one:** $A^+ \equiv AA^*$
- **Union:** $A \mid B \equiv A+B$
- **Option:** $A? \equiv A + \varepsilon$
- **Range:** $'a' + 'b' + \dots + 'z' \equiv [a-z]$
- **Excluded range:** complement of $[a-z] \equiv [^a-z]$

RE Examples

Regular Expression	Explanation
a^*	0 or more a's (ϵ , a, aa, aaa, aaaa, ...)
a^+	1 or more a's (a, aa, aaa, aaaa, ...)
$(a b)(a b)$	(aa, ab, ba, bb)
$(a b)^*$	all strings of a's and b's (including ϵ)
$(aa ab ba bb)^*$	all strings of a's and b's of even length
$[a-zA-Z]$	shorthand for " $a b \dots z A B \dots Z$ "
$[0-9]$	shorthand for " $0 1 2 \dots 9$ "
$0([0-9])^*0$	numbers that start and end with 0
$1^*(0 \epsilon)1^*$	binary strings that contain at most one zero
$(0 1)^*00(0 1)^*$	all binary strings that contain '00' as substring

- Q: are $(a|b)^*$ and $(a^*b^*)^*$ equivalent?

More Examples

- Keywords: 'if' or 'else' or 'then' or 'for' ...
 - RE = 'i''f' + 'e''l''s''e' + ... = 'if' + 'else' + 'then' + ...
- Numbers: a non-empty string of digits
 - digit = '0' + '1' + '2' + '3' + '4' + '5' + '6' + '7' + '8' + '9'
 - integer = digit digit*
 - Q: is '000' an integer?
- Identifier: strings of letters or digits, starting with a letter
 - letter = 'a' + 'b' + ... 'z' + 'A' + 'B' + ... + 'Z' = [a-zA-Z]
 - RE = letter(letter + digit)*
 - Q: is the RE valid for identifiers in C?
- Whitespace: a non-empty sequence of blanks, newline and tabs
 - (' ' + '\n' + '\t')+

REs in Programming Language

Symbol	Meaning		
\d	Any decimal digit, i.e. [0-9]		
\D	Any non-digit char, i.e., [^0-9]		
\s	Any whitespace char, i.e., [\t\n\r\f\v]		
\S	Any non-whitespace char, i.e., [^ \t\n\r\f\v]		
\w	Any alphanumeric char, i.e., [a-zA-Z0-9_]		
\W	Any non-alphanumeric char, i.e., [^a-zA-Z0-9_]		
.	Any char	\.	Matching “.”
[a-f]	Char range	[^a-f]	Exclude range
^	Matching string start	\$	Matching string end
(...)	Capture matches		

<https://docs.python.org/3/howto/regex.html>

Lexical Specification of a Language

- S0: write a regex for the lexemes of each token class
 - Numbers = digit^+
 - Keywords = 'if' + 'else' + ...
 - Identifiers = $\text{letter}(\text{letter} + \text{digit})^*$
- S1: construct R, matching all lexemes for all tokens
 - $R = \text{numbers} + \text{keywords} + \text{identifiers} + \dots = R1 + R2 + R3 + \dots$
- S2: let input be $x_1 \dots x_n$, for $1 \leq i \leq n$, check $x_1 \dots x_i \in L(R)$
- S3: if successful, then we know $x_1 \dots x_i \in L(R_j)$ for some j
- S4: remove $x_1 \dots x_i$ from input and go to step S2

Lexical Specification of a Language

- How much input is used?
 - $x_1 \dots x_i \in L(R), x_1 \dots x_j \in L(R), i \neq j$
 - Which one do we want? (e.g., '==' or '=')
 - Maximal match: always choose the longer one
- Which token is used if more than one matches?
 - $x_1 \dots x_i \in L(R)$ where $R = R_1 + R_2 + \dots + R_n$
 - $x_1 \dots x_i \in L(R_m), x_1 \dots x_i \in L(R_n), m \neq n$
 - E.g., keywords = 'if', identifier = letter(letter+digit)*
 - Keyword has higher priority
 - Rule of thumb: choose the one listed first
- What if no rule matches?
 - $x_1 \dots x_i \notin L(R) \rightarrow \text{Error}$

Summary: RE

- We have learnt how to specify tokens for lexical analysis
 - Regular expressions
 - Concise notations for the string patterns
- Used in lexical analysis with some extensions
 - To resolve ambiguities
 - To handle errors
- REs is only a language specification
 - An implementation is still needed
 - Next: to construct a token recognizer for languages given by regular expressions – by using **finite automata**